

Class 9: Candy Mini-Project

Seunghoon Oh (PID: A19372132)

Table of contents

Background	1
Importing candy data	1
What is in the dataset?	2
What is your favorite candy?	3
Side-note: the skimr::skim() function	4
Exploratory analysis	5
Overall Candy Rankings	9
Time to add some useful color	15
Taking a look at pricepercent	17
Exploring the correlation structure	19
Principal Component Analysis	21
Using plotly to generate interactive plots	25
Summary	26

Background

Today we are taking a detour to analyze a fun dataset (that we have more intrinsic insight into) with the most useful analysis method we have learned thus far - Principal Component Analysis (PCA)

The data is all related to halloween candy and is from the 538 website.

Importing candy data

```
candy_file <- "candy-data.csv"

candy = read.csv(candy_file, row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer
100 Grand	1	0	1	0	0	1
3 Musketeers	1	0	0	0	1	0
One dime	0	0	0	0	0	0
One quarter	0	0	0	0	0	0
Air Heads	0	1	0	0	0	0
Almond Joy	1	0	0	1	0	0
	hard bar	pluribus	sugarpercent	pricepercent	winpercent	
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

What is in the dataset?

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

```
[1] 85
```

=> there are 85 rows (which indicate different candy types) in this dataset.

Q2. How many fruity candy types are in the dataset?

```
nrow(candy[candy$fruity == 1, ])
```

```
[1] 38
```

```
table(candy$fruity)
```

```
0 1
47 38
```

=> There are 38 fruity candy types in the dataset.

What is your favorite candy?

```
candy["Twix", ]$winpercent
```

```
[1] 81.64291
```

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
candy |>
  filter(row.names(candy)=="Twix") |>
  select(winpercent)
```

```
      winpercent
Twix    81.64291
```

Q3. What is your favorite candy (other than Twix) in the dataset and what is it's winpercent value?

```
candy["Junior Mints", ]$winpercent
```

```
[1] 57.21925
```

=> 57.21925

Q4. What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat", ]$winpercent
```

```
[1] 76.7686
```

```
=> 76.7686
```

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

```
=> 49.6535
```

Side-note: the `skimr::skim()` function

```
#library("skimr")  
skimr::skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_vari- able	n_miss- ing	com- plete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	

skim_vari- able	n_miss- ing	com- plete_rate	mean	sd	p0	p25	p50	p75	p100	hist
peanutyal- mondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedrice- wafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

9 variables, from **chocolate** to **pluribus**, has a value of either 0 or 1. (**int** type) Other 3 variables that end with **percent** is a **numeric** type value. These store an actual number rather than categorical assignment as other 9 variables.

And among those 3 variables, **winpercent** has a different scale because values for other “percent” variables are between 0 and 1, while this variable has values between 0 and 100.

Q7. What do you think a zero and one represent for the `candy$chocolate` column?

I think it is basically same as a logical value -> 0 represents **False** and 1 represents **True**. So for the `candy$chocolate` column, 0 would indicate that the candy type is **not a chocolate** (**type candy**) and 1 would indicate that it **is a chocolate** (**type candy**)


Exploratory analysis

Q8. Plot a histogram of `winpercent` values using both base R and `ggplot2`.

Using base R:

```
hist(candy$winpercent)
```


Histogram of candy\$winpercent



Using ggplot2:

```
library(ggplot2)
ggplot (candy, aes(winpercent)) +
  geom_histogram()
```

``stat_bin()`` using ``bins = 30``. Pick better value ``binwidth``.



Q9. Is the distribution of winpercent values symmetrical?

From the plots in Q8, **No**.

Q10. Is the center of the distribution above or below 50%?

Center of the distribution would be:

```
median(candy$winpercent)
```

```
[1] 47.82975
```

which is **below** 50%.

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
chocolate_winpercent <- mean(candy$winpercent[candy$chocolate == 1])
fruity_winpercent <- mean(candy$winpercent[candy$fruity == 1])
cat ("chocolate_winpercent", chocolate_winpercent, "\n")
```

```
chocolate_winpercent 60.92153
```

```
cat ("fruity_winpercent", fruity_winpercent, "\n")
```

```
fruity_winpercent 44.11974
```

```
chocolate_winpercent > fruity_winpercent
```

```
[1] TRUE
```

=> Chocolate candies have **higher** winpercent average than fruity candies.

Using `as.logical()`:

```
mean(candy$winpercent[as.logical(candy$chocolate)])
```

```
[1] 60.92153
```

```
mean(candy$winpercent[as.logical(candy$fruity)])
```

```
[1] 44.11974
```

Q12. Is this difference statistically significant?

```
choc_winperc_data <- candy$winpercent[as.logical(candy$chocolate)]
fruit_winperc_data <- candy$winpercent[as.logical(candy$fruity)]

t.test (choc_winperc_data, fruit_winperc_data)
```

Welch Two Sample t-test

```
data:  choc_winperc_data and fruit_winperc_data
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153  44.11974
```

=> The difference is statistically significant, as it has a very low p-value ($\ll 0.05$). This indicates that the null hypothesis is rejected and that the observed difference is likely not due to chance.

Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

Using base R:

```
head(candy[order(candy$winpercent),], n=5)
```

	chocolate	fruity	caramel	peanut	almond	nougat		
Nik L Nip	0	1	0		0	0		
Boston Baked Beans	0	0	0		1	0		
Chiclets	0	1	0		0	0		
Super Bubble	0	1	0		0	0		
Jawbusters	0	1	0		0	0		

	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	price	percent
Nik L Nip			0	0	0	1		0.197		0.976
Boston Baked Beans			0	0	0	1		0.313		0.511
Chiclets			0	0	0	1		0.046		0.325
Super Bubble			0	0	0	0		0.162		0.116
Jawbusters			0	1	0	1		0.093		0.511

	winpercent
Nik L Nip	22.44534
Boston Baked Beans	23.41782
Chiclets	24.52499
Super Bubble	27.30386
Jawbusters	28.12744

Using dplyr:

```
candy |>
  arrange(winpercent) |>
  head(5)
```

	chocolate	fruity	caramel	peanut	almond	nougat		
Nik L Nip	0	1	0		0	0		
Boston Baked Beans	0	0	0		1	0		
Chiclets	0	1	0		0	0		
Super Bubble	0	1	0		0	0		
Jawbusters	0	1	0		0	0		

	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	price	percent
Nik L Nip			0	0	0	1		0.197		0.976

Boston Baked Beans	0	0	0	1	0.313	0.511
Chiclets	0	0	0	1	0.046	0.325
Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511

	winpercent
Nik L Nip	22.44534
Boston Baked Beans	23.41782
Chiclets	24.52499
Super Bubble	27.30386
Jawbusters	28.12744

=> Nik L Nip, Boston Baked Beans, Chiclets, Super Bubble, Jawbusters

Q14. What are the top 5 all time favorite candy types out of this set?

```
head(candy[order(candy$winpercent, decreasing=TRUE),], n=5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1

	crisped	rice	wafer	hard	bar	pluribus	sugar
Reese's Peanut Butter cup		0	0	0		0	0.720
Reese's Miniatures		0	0	0		0	0.034
Twix		1	0	1		0	0.546
Kit Kat		1	0	1		0	0.313
Snickers		0	0	1		0	0.546

	price	percent	winpercent
Reese's Peanut Butter cup	0.651		84.18029
Reese's Miniatures	0.279		81.86626
Twix	0.906		81.64291
Kit Kat	0.511		76.76860
Snickers	0.651		76.67378

```
candy |>
  arrange(desc(winpercent)) |>
  head(5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1

	crisped	rice	wafer	hard	bar	pluribus	sugar	percent
Reese's Peanut Butter cup		0	0	0		0		0.720
Reese's Miniatures		0	0	0		0		0.034
Twix		1	0	1		0		0.546
Kit Kat		1	0	1		0		0.313
Snickers		0	0	1		0		0.546


	price	percent	win	percent
Reese's Peanut Butter cup	0.651		84.18029	
Reese's Miniatures	0.279		81.86626	
Twix	0.906		81.64291	
Kit Kat	0.511		76.76860	
Snickers	0.651		76.67378	

=> **Reese's Peanut Butter cup, Reese's Miniatures, Twix, Kit Kat, Snickers**

Q15. Make a first barplot of candy ranking based on winpercent values.

```
library(ggplot2)


ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```



Q16. This is quite ugly, use the `reorder()` function to get the bars sorted by `winpercent`?

```
library(ggplot2)


ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col()
```



Time to add some useful color

```
my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "red"

ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=my_cols) +
  ylab("")
```



Q17. What is the worst ranked chocolate candy?

=> Sixlets


Q18. What is the best ranked fruity candy?

=> Starburst

Taking a look at pricepercent

```
library(ggrepel)

# How about a plot of win vs price
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

```
ord <- order(candy$winpercent/candy$pricepercent, decreasing = TRUE)
head( candy[ord,c(11,12)], n=5 )
```

	pricepercent	winpercent
Tootsie Roll Midgies	0.011	45.73675
Pixie Sticks	0.023	37.72234
Fruit Chews	0.034	43.08892
Dum Dums	0.034	39.46056
Strawberry bon bons	0.058	34.57899

=> I sorted the candy types by winpercent/pricepercent value. The candy type with highest value is **Tootsie Roll Midgies**.

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

Finding the top 5 most expensive candy types

Using base R:

```
ord <- order(candy$pricepercent, decreasing = TRUE)
head( candy[ord,c(11,12)], n=5 )
```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Using dplyr:

```
library(dplyr)
candy |>
  arrange(-pricepercent) |>
  select(pricepercent, winpercent) |>
  head(n=5)
```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

=> the top 5 most expensive candy types are: **Nik L Nip**, **Nestle Smarties**, **Ring pop**, **Hershey's Krackel**, **Hershey's Milk Chocolate**.

=> from the table, we can say that **Nik L Nip** is the least popular among these five candy types.

Exploring the correlation structure


From the sidenote: This is also why we will use `scale=TRUE` when calling `prcomp()` in the next section: it ensures all variables contribute equally by working from correlations (standardized and scale-free) rather than raw covariances, which would be dominated by whichever variable happens to have the largest numeric range. Can you guess which variable in our dataset that would be?

The variable that would dominate the `prcomp()` result without `scale=TRUE` argument setting would be the `winpercent` variable, as it has the largest numeric range in the dataset.

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
cij <- cor(candy)
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

=> Some of the most evident anti-correlated pairs from this figure is: **chocolate-fruity**, **bar-fruity**, **bar-pluribus** and **fruity-pricepercent**, **fruity-winpercent**, **fruity-peanutyalmondy**, etc.

Q23. Use your corrpilot result to make a prediction: which variables do you expect will have the largest contributions (i.e. loadings) to PC1 (i.e., drive the most separation between candies along the first principal component)?

fruity and chocolate. They have the largest magnitude of negative correlation, which shows that they would be likely to be the most “separated” variables in the dataset.

Principal Component Analysis

In this case we want to be sure to set `scale=TRUE` argument for `prcomp()` because we have one variable `winpercent`, that is on a very different scale than all others and would otherwise dominate our PCA results.


```
pca <- prcomp(candy, scale=TRUE)
summary(pca)
```

Importance of components:


	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369

	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000

```
plot(pca$x[, c("PC1", "PC2")])
```




```
plot(pca$x[,1:2], col=my_cols, pch=16)
```



```
# Make a new data-frame with our PCA results and candy data
my_data <- cbind(candy, pca$x[,1:3])

p <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=my_cols)

p
```




```
library(ggrepel)

p + geom_text_repel(size=3.3, col=my_cols, max.overlaps = 7) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
        subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown),",
        caption="Data from 538")
```


Halloween Candy PCA Space

Colored by type: chocolate bar (dark brown), chocolate other (light brown), fruity (red), other (black)



Data from 538


Using plotly to generate interactive plots

```
#library(plotly)
#ggplotly(p)
```

Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to

you? Where did you see this relationship highlighted previously?

```
ggplot(pca$rotation) +  
  aes(PC1, reorder(colnames(candy), PC1)) +  
  geom_col()
```



=> **fruity** has the most positive PC1 value among the variables. Other variables with positive PC1 value include **pluribus** and **hard**. Comparing this result with the previously generated **corrplot** *makes sense*, since **fruity** has shown positive correlation with **pluribus** and **hard**, while showing significant anti-correlation with the top negative variables in terms of PC1 - **chocolate**, **bar**, **winpercent**, **pricepercent** or **peanutyalmondy**. So the PCA captures large proportion of these complex relationships between variables in this axis PC1.

Summary

Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?

Making a “winning” candy - an **expensive peanutyalmondy chocolate bar**. characteristics **fruity**, **hard**, or **pluribus** should preferably not be included here. Other attributes - **caramel**, **nougat**, **crispedricewafer** or **sugarpercent** - would have a positive effect (positive correlation) in

making a winning candy, but the effects are much more moderate than the top 4 attributes indicated as bold.